

## CS6560: Parallel Computer Architecture

### Basics of Parallel Programming



Madhu Mutyam  
PACE Laboratory  
Department of Computer Science and Engineering  
Indian Institute of Technology Madras



Feb 5 - 12, 2018

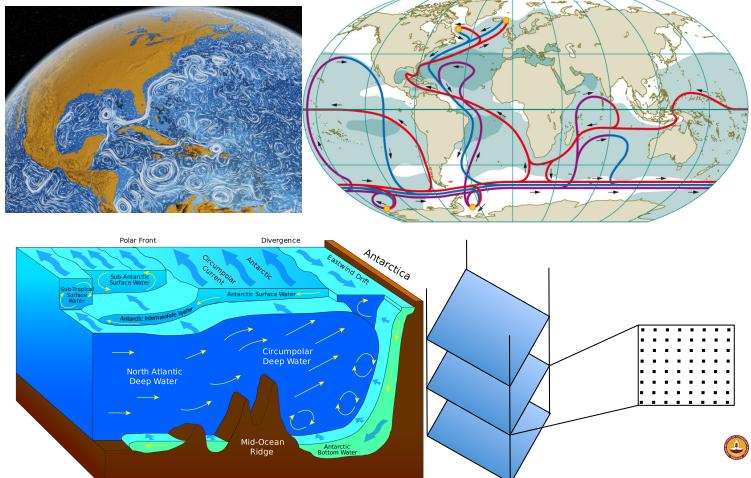
## Understand Program Behavior

- ▶ Understand the behavior of programs that run on a computing system
  - ▶ Principle of locality – Design of memory hierarchy
  - ▶ Type of operations – Design of ISA
  - ▶ Frequency of various operations – Efficient design of common case
- ▶ Important for uniprocessor as well as multiprocessor systems
- ▶ Important for algorithm designers, programmers, and architects

## Parallel Application Case Studies

- ▶ Simulating Ocean Currents
- ▶ Simulating Evolution of Galaxies
- ▶ Visualizing Complex Scenes using Ray Tracing
- ▶ Mining Data for Associations

## Simulating Ocean Currents



Madhu Mutyam (IIT Madras)

Feb 5 - 12, 2018

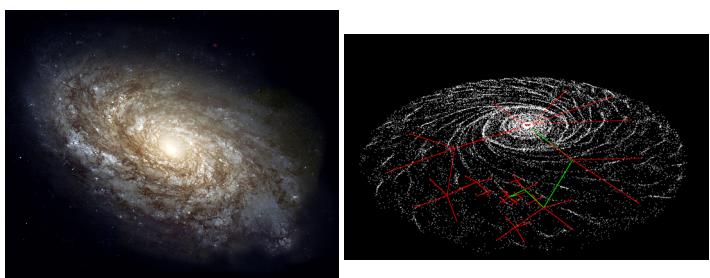
2/32

Madhu Mutyam (IIT Madras)

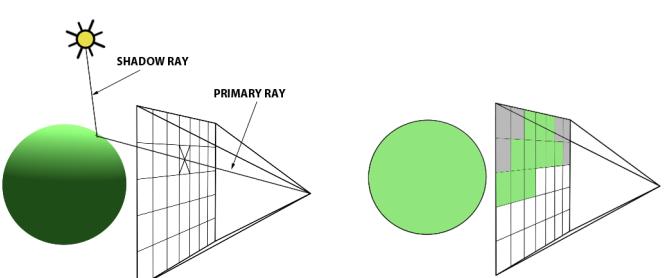
Feb 5 - 12, 2018

3/32

## Simulating Evolution of Galaxies



## Visualizing Complex Scenes using Ray Tracing



Madhu Mutyam (IIT Madras)

Feb 5 - 12, 2018

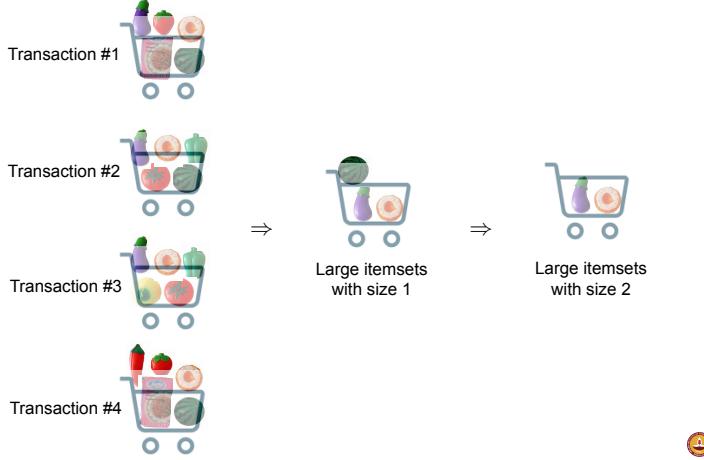
4/32

Madhu Mutyam (IIT Madras)

Feb 5 - 12, 2018

5/32

## Mining Data for Associations

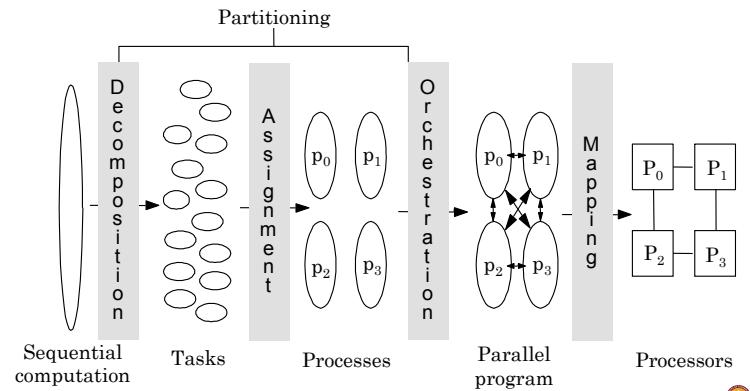


Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

6/32

## The Parallelization Process



Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

7/32

## Decomposition

- ▶ Breaking the computation into a collection of tasks
- ▶ Provides an upper bound on the number of processes
- ▶ Expose enough concurrency without having too much overhead in managing the tasks

Madhu Mutyan (IIT Madras)

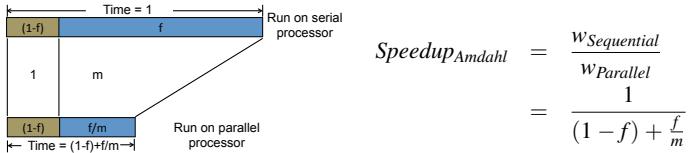
Feb 5 - 12, 2018

8/32

## Amdahl's Law<sup>a</sup>

<sup>a</sup>Amdahl. Validity of the Single-Processor Approach to Achieving Large-Scale Computing Capabilities. AFIPS, 1967.

- ▶ Consider a workload  $w$ , where fraction  $f$  of work can be executed in parallel and the remaining in serial



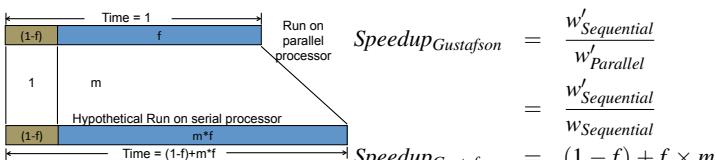
$$\begin{aligned} \text{Speedup}_{\text{Amdahl}} &= \frac{w_{\text{Sequential}}}{w_{\text{Parallel}}} \\ &= \frac{1}{(1-f) + \frac{f}{m}} \end{aligned}$$

- ▶ For small values of  $f$ , the enhancements will have little effect
- ▶ As  $m \rightarrow \infty$ , the speedup is bound by  $\frac{1}{(1-f)}$
- ▶ Favours single processors
- ▶ Fixed-size speedup model

## Gustafson's Law<sup>a</sup>

<sup>a</sup>Gustafson. Reevaluating Amdahl's Law. Comm. ACM, 1988.

- ▶ A machine with parallel computation ability lets more computations in the same amount of time
- ▶  $w$  and  $w' = ((1-f) + f * m) * w$  take the same amount of time executing with serial processing and parallel processing with  $m$  processors, resp.



$$\begin{aligned} \text{Speedup}_{\text{Gustafson}} &= \frac{w'_{\text{Sequential}}}{w'_{\text{Parallel}}} \\ &= \frac{w'_{\text{Sequential}}}{w_{\text{Sequential}}} \\ &= (1-f) + f \times m \end{aligned}$$

- ▶ The speedup is a linear function of  $m$  if the workload is scaled up to maintain a fixed execution time
- ▶ Favours large-scale parallel system
- ▶ Fixed-time speedup model

Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

10/32

## Amdahl's Law in the Multicore Era<sup>a</sup>

<sup>a</sup>Hill and Marty. Amdahl's Law in the Multicore Era. IEEE Computer, 2008.

- ▶ Resources in a multicore system are bounded, say  $n$
- ▶ Microarchitects can improve single-core performance using more of the bounded resources
- ▶ A simple baseline core consumes 1 Base Core Equivalent (BCE) resources and provides normalized performance of 1
- ▶ An enhanced core consumes  $r$  BCE resources and provides normalized performance of  $\text{perf}(r)$ 
  - ▶ if  $\text{perf}(r) > r$ , speeds up for both sequential and parallel execution
  - ▶ if  $\text{perf}(r) < r$ , benefits sequential execution, but hurts parallel execution
  - ▶  $\text{perf}(r) = \sqrt{r}$  (Pollack's Law)

Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

11/32

## Symmetric Multicore Chips<sup>a</sup>

<sup>a</sup>Hill and Marty. Amdahl's Law in the Multicore Era., IEEE Computer, 2008.

- ▶ The number of resources in a multicore chip is  $n$
- ▶ The number of cores is  $\frac{n}{r}$  with  $r$  BCE resources per core
- ▶ Serial fraction  $(1-f)$  uses 1 core at rate  $perf(r)$
- ▶ Serial time =  $\frac{1-f}{perf(r)}$
- ▶ Parallel fraction ( $f$ ) uses  $\frac{n}{r}$  cores at rate  $perf(r)$  each
- ▶ Parallel time =  $\frac{f}{perf(r) \times \frac{n}{r}} = \frac{f \times r}{perf(r) \times n}$
- ▶ Speedup of symmetric multicore w.r.t. one base core:
 
$$Speedup_{Symmetric} = \frac{1}{\frac{(1-f)}{perf(r)} + \frac{f \times r}{perf(r) \times n}}$$
- ▶ If  $m = \frac{n}{r}$ ,
 
$$Speedup_{Symmetric} = \frac{perf(r)}{(1-f) + \frac{f}{m}}$$
- ▶ Amdahl's law, fixed-size speedup model for multicore

$$Speedup_{Symmetric} = \frac{1}{\frac{(1-f)}{perf(r)} + \frac{f \times r}{perf(r) \times n}}$$

$$Speedup_{Symmetric} = \frac{perf(r)}{(1-f) + \frac{f}{m}}$$



## Asymmetric Multicore Chips<sup>a</sup>

<sup>a</sup>Hill and Marty. Amdahl's Law in the Multicore Era. IEEE Computer, 2008.

- ▶ The number of resources in a multicore chip is  $n$
- ▶ Enhance some (but not all) cores
- ▶ One enhanced core ( $r$ -BCE) and  $(n-r)$  number of 1-BCE cores
- ▶ The number of cores is  $1 + n - r$
- ▶ Serial fraction  $(1-f)$  uses 1 enhanced core at rate  $perf(r)$
- ▶ Serial time =  $\frac{1-f}{perf(r)}$
- ▶ Parallel fraction ( $f$ ) uses one core at  $perf(r)$  and  $(n-r)$  cores at rate 1
- ▶ Parallel time =  $\frac{f}{perf(r)+n-r}$
- ▶ Speedup of asymmetric multicore w.r.t. one base core:

$$Speedup_{Asymmetric} = \frac{1}{\frac{(1-f)}{perf(r)} + \frac{f}{perf(r)+n-r}}$$



## Fixed-Time Speedup Model for Symmetric Multicore<sup>a</sup>

<sup>a</sup>Sun and Chen. Reevaluating Amdahl's Law in the Multicore Era. JPDC, 2010.

- ▶ From the fixed-size speedup model for symmetric multicore, the speedup achieved with  $\frac{n}{r}$  cores, each having  $r$  BCE resources, is  $\frac{1}{\frac{(1-f)}{perf(r)} + \frac{f \times r}{perf(r) \times n}}$
- ▶ If  $n = r$ , the speedup achieved with one core is  $perf(r)$
- ▶ Time ( $t$ ) to execute work  $w$  using one core at  $perf(r)$  is  $\frac{(1-f) \times w}{perf(r)} + \frac{f \times w}{perf(r)}$
- ▶  $w'$  is the amount of work carried out by a multicore system with  $m$  cores such that  $t = \frac{(1-f) \times w}{perf(r)} + \frac{f \times w'}{perf(r) \times m}$
- ▶ Thus,  $w' = w \times m$
- ▶ The scaled up speedup, w.r.t. to single core system, is

$$\begin{aligned} Speedup_{FT\_Symmetric} &= \frac{w'_{Sequential}}{w_{Sequential}} = \frac{\frac{(1-f) \times w}{perf(r)} + \frac{f \times w'}{perf(r)}}{\frac{w}{perf(r)}} \\ &= (1-f) + m \times f \end{aligned}$$

## Assignment

- ▶ Mechanism by which tasks will be distributed among processes
- ▶ Inspection of the code or high-level understanding of the application can help in appropriate assignment
- ▶ Static vs dynamic assignment
- ▶ Performance goals:
  - ▶ Workload balancing
  - ▶ Reducing interprocess communication
  - ▶ Reducing run-time overhead of managing assignment
- ▶ Decomposition and assignment are the major algorithmic steps in parallelization
  - ▶ Independent of the underlying architecture and programming model



## Orchestration

- ▶ Deals with naming, organizing data structures, scheduling tasks, communicating, synchronizing
- ▶ Performance goals:
  - ▶ Reducing the cost of communication and synchronization
  - ▶ Preserving data locality
  - ▶ Scheduling tasks to satisfy dependences early
  - ▶ Reducing the overhead of parallelism management
- ▶ Architecture, programming model, and programming language play a big role

## Mapping

- ▶ Deals with which process to run on which processor
- ▶ Specific to the system or programming environment
- ▶ Processes can be pinned to processors or controlled by OS dynamically
- ▶ Performance goals:
  - ▶ Keeping related processes in the same processor, if necessary
  - ▶ Exploiting locality in network topology



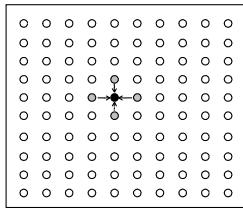
## The Gauss-Seidel Equation Solver Kernel

```

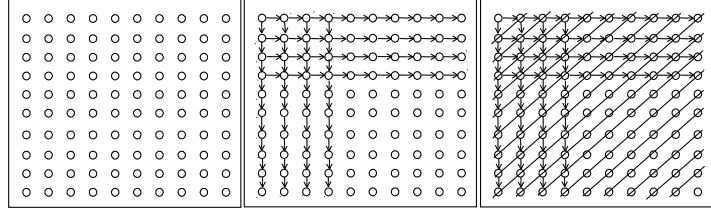
1. int n;
2. float **A, diff=0;
3. main()
4. begin
5.   read(n);
6.   A← malloc(a 2-d array of size n+2-by-n+2 doubles);
7.   Initialize (A);
8.   Solve (A);
9. end main

10. procedure Solve (A)
11. float **A;
12. begin
13.   int i, j, done=0;
14.   float diff=0, temp;
15.   while (!done) do
16.     diff = 0;
17.     for i ← 1 to n do
18.       for j ← 1 to n do
19.         temp = A[i,j];
20.         A[i,j] ← 0.2*(A[i,j]+A[i,j-1]+A[i-1,j]+A[i,j+1]+A[i+1,j]);
21.         diff += abs(A[i,j]-temp);
22.       end for
23.     end for
24.     if (diff/(n*n)<TOL) then done = 1;
25.   end while
26. end procedure

```



## Dependences and Concurrency in the Gauss-Seidel Equation Solver Kernel



- ▶ Out of  $O(n^2)$  work involved in each sweep, there exist  $O(n)$  concurrency and  $O(n)$  sequential dependence along anti-diagonals



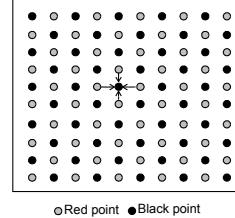
## Exploiting Concurrency in the Gauss-Seidel Equation Solver Kernel

- ▶ Assume the computation is decomposed at the grid-point granularity
- ▶ Retain the loop structure and insert point-to-point synchronization
  - ▶ the new value for a grid point has to be produced in the current sweep before it is used by the points below and to its right
  - ▶ Synchronization overhead is too high
- ▶ Restructure the loops and use global synchronization between iterations
  - ▶ The outer *for* loop can be over anti-diagonals and the inner *for* loop can be over elements within an anti-diagonal
  - ▶  $O(n)$  synchronizations and load imbalance



## Exploiting Knowledge of Problem for Efficient Solution

- ▶ As Gauss-Seidel method is not an exact solution method, we can update grid points in *red-black* ordering



○ Red point ● Black point

- ▶ Compute all  $(\frac{n^2}{2})$  red points in parallel ⇒ Global synchronization ⇒ compute all  $(\frac{n^2}{2})$  black points in parallel
- ▶ Produced results are independent of the number of processors



## Decomposition

- ▶ Decompose into individual inner loop iterations
  - ▶ Degree of concurrency is  $n^2$

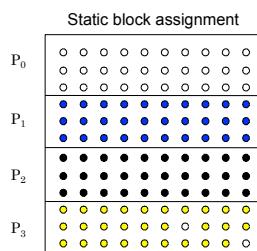
```

15. while (!done) do      /* a sequential loop */
16.   diff = 0;
17.   for_all i ← 1 to n do /* a parallel loop nest */
18.     for_all j ← 1 to n do
19.       temp = A[i,j];
20.       A[i,j] ← 0.2*(A[i,j]+A[i,j-1]+A[i-1,j]+A[i,j+1]+A[i+1,j]);
21.       diff += abs(A[i,j]-temp);
22.     end for_all
23.   end for_all
24.   if (diff/(n*n) < TOL) then done = 1;
25. end while

```

## Assignment

- ▶ Assume row-based decomposition;  $n$  rows and  $p$  processes
- ▶ Static Assignment
  - ▶ Cyclic assignment: process  $i$  gets rows  $i, i + p, \dots$
  - ▶ Block assignment: each process gets a group of rows



$$\frac{\text{Communication}}{\text{Computation}} = O(\frac{p}{n})$$

- ▶ Dynamic Assignment

- ▶ Each process gets a next available row



## Orchestration Under the Data Parallel Model

```

1. int n, nprocs;
2. float **A, diff=0;
3. main()
4. begin
5.   read(n); read(nprocs);
6.   A ← G_MALLOC(2-d array of
      size(n+2)×(n+2) doubles);
7.   initialize(A);
8.   Solve(A);
9. end main

10. procedure Solve(A)
11.   float **A;
12. begin
13.   int i, j, done = 0;
14.   float temp, mydiff = 0;
14a.  DECOMP A[BLOCK, *, nprocs];
15.   while (!done) do
16.     mydiff = 0;
17.     for_all i ← 1 to n do
18.       for_all j ← 1 to n do
19.         temp = A[i, j];
20.         A[i, j] = 0.2*(A[i, j] + A[i, j-1] + A[i-1, j]
                         + A[i, j+1] + A[i+1, j]);
21.         mydiff += abs(A[i, j]-temp);
22.     end for_all
23.   end for_all
24a.  REDUCE(mydiff, diff, ADD);
25.   if(diff/(n*n) < TOL) then done = 1;
26. end while
27. end procedure

```

Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

34/32

## Orchestration Under the Shared Memory Model

```

1. int n, nprocs;
2a. float **A, diff;
2b. LOCKDEC(diff_lock);
2c. BARDEC (bar1);

3. main()
4. begin
5.   read(n); read(nprocs);
6.   A ← G_MALLOC(2-d array of size(n+2)×(n+2) doubles);
7.   initialize(A);
8a. CREATE(nprocs-1, Solve, A);
8.   Solve(A);
8b. WAIT_FOR_END(nprocs-1);
9. end main

```

Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

25/32

## Orchestration Under the Shared Memory Model (Contd)

```

10. procedure Solve(A)
11.   float **A;
12. begin
13.   int i, j, pid, done = 0;
14.   float temp, mydiff = 0;
14a.  int mymin = 1 + (pid * n/nprocs);
14b.  int mymax = mymin + n/nprocs - 1;
15.   while (!done) do
16.     mydiff = diff = 0;
16a.  BARRIER(bar1, nprocs);
17.     for i ← mymin to mymax do
18.       for j ← 1 to n do
19.         temp = A[i, j];
20.         A[i, j] = 0.2*(A[i, j] + A[i, j-1] + A[i-1, j]
                         + A[i, j+1] + A[i+1, j]);
21.         mydiff += abs(A[i, j]-temp);
22.     end for
23.   end for
24a.  LOCK(diff_lock);
24b.  diff += mydiff;
24c.  UNLOCK(diff_lock);
24d.  BARRIER(bar1, nprocs);
24e.  if (diff/(n*n) < TOL) then done = 1;
24f.  BARRIER(bar1, nprocs);
25. end while
26. end procedure

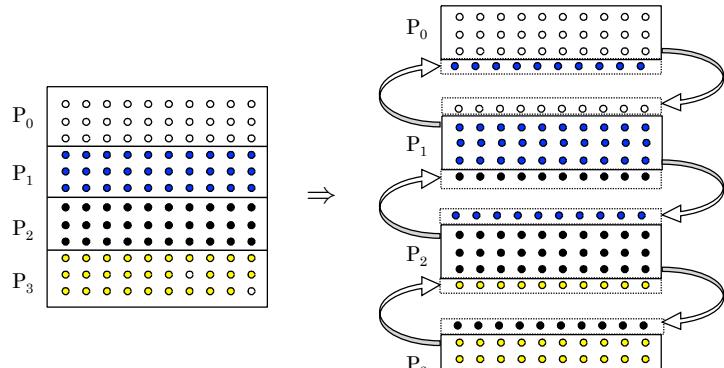
```

Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

26/32

## Data Layout Under the Message Passing Model



Feb 5 - 12, 2018

27/32

## Orchestration Under the Message Passing Model

```

1. int n, pid, nprocs;
2. float **myA;
3. main()
4. begin
5.   read(n); read(nprocs);
8a.  CREATE(nprocs-1, Solve);
8b.  Solve();
8c.  WAIT_FOR_END(nprocs-1);
9. end main

10. procedure Solve()
11. begin
12.   int i, j, pid, n'=n/nprocs, done=0;
13.   float temp, tempdiff, mydiff=0;
6.   myA ← malloc(2-d array of size(n/nprocs+2)×(n+2));
7.   initialize(myA);
15.   while(!done) do
16.     mydiff = 0;
16a.  if(pid!=0) SEND(&myA[1,0], n*sizeof(float), pid-1, ROW);
16b.  if(pid!=nprocs-1) SEND(&myA[n',0], n*sizeof(float),
                           pid+1, ROW);
16c.  if(pid!=0) RECEIVE(&myA[0,0], n*sizeof(float),
                           pid-1, ROW);
16d.  if(pid!=nprocs-1) RECEIVE(&myA[n'+1,0], n*sizeof(float),
                           pid+1, ROW);

```

Madhu Mutyan (IIT Madras)

Feb 5 - 12, 2018

28/32

## Orchestration Under the Message Passing Model (Contd)

```

17.   for i ← 1 to n' do
18.     for j ← 1 to n do
19.       temp = myA[i, j];
20.       myA[i, j]=0.2*(myA[i, j] + myA[i, j-1] +
                         myA[i-1, j] + myA[i, j+1] + myA[i+1, j]);
21.       mydiff += abs(A[i, j]-temp);
22.     end for
23.   end for
24a.  if(pid!=0) then
24b.    SEND(mydiff, sizeof(float), 0, DIFF);
24c.    RECEIVE(done, sizeof(int), 0, DONE);
24d.  else
24e.    for i ← 1 to nprocs-1 do
24f.      RECEIVE(tempdiff, sizeof(float), *, DIFF);
24g.      mydiff += tempdiff;
24h.    end for
24i.    if(mydiff/(n*n) < TOL) then done = 1;
24j.    for i ← 1 to nprocs-1 do
24k.      SEND(done, sizeof(int), i, DONE);
24l.    end for
24m.  end if
25. end while
26. end procedure

```

Madhu Mutyan (IIT Madras)

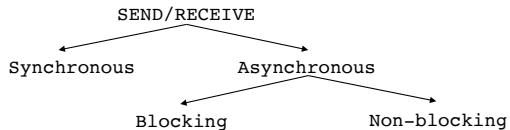
Feb 5 - 12, 2018

29/32

## SEND/RECEIVE Alternatives

- ▶ Differ in their completion semantics – based on when they return control back to issuing process

- ▶ Affect when the data structures/buffers can be reused without violating correctness



- ▶ Synchronous messages provide built-in synchronization through match
  - ▶ Separate event synchronizations are required for asynchronous messages

## Summary

- ▶ Parallelizing a sequential application:

- ▶ Depose the work into tasks
- ▶ Assign the tasks to processes
- ▶ Orchestrate data accesses, communication, and synchronization among processes
- ▶ Map processes to processors

- ▶ Decomposition and assignment are independent of programming models and underlying hardware

- ▶ Communication is implicit in shared memory (SM) model while it is explicit in message passing (MP) model

- ▶ Synchronization primitives are required in SM but not in MP

# Thank You

